

Manipulating Data in R

Recap of Data Cleaning

- `is.na()`, `any(is.na())`, `all(is.na())`, `count()`, and functions from `naniar` like `gg_miss_var()` and `miss_var_summary` can help determine if we have NA values
- `miss_var_which()` can help you drop columns that have any missing values.
- `filter()` automatically removes NA values
- `drop_na()` can help you remove NA values
- NA values can change your calculation results
- think about what NA values represent - don't drop them if you shouldn't
- `replace_na()` will replace `NA values with a particular value

Recap of Data Cleaning

- `case_when()` can recode **entire values** based on conditions
 - remember `case_when()` needs `TRUE ~ variable` to keep values that aren't specified by conditions, otherwise will be `NA`
- `stringr` package has great functions for looking for specific **parts of values** especially `filter()` and `str_detect()` combined
 - also has other useful string manipulation functions like `str_replace()` and more!
 - `separate()` can split columns into additional columns
 - `unite()` can combine columns

□ [Cheatsheet](#)

Manipulating Data

In this module, we will show you how to:

1. Reshape data from wide to long
2. Reshape data from long to wide
3. Merge Data/Joins

What is wide/long data?

Data is wide or long **with respect to** certain variables.

The diagram illustrates the transformation of data from a wide format to a long format. On the left, a wide data table is shown with two rows for Patient 1 and Patient 2. Each row has three columns for Day 1, Day 2, and Day 3, with values A, B, C for Patient 1 and D, E, F for Patient 2. An arrow points from this wide table to the right, labeled "Wide" above the arrow and "Long" below it. On the right, a long data table is shown with six rows. The first three rows represent Patient 1, and the last three represent Patient 2. Each row has two columns: "Day" and "Value". The data is repeated across rows for each day, showing the "long" structure where multiple observations are grouped by a common identifier (Patient ID) and time point (Day).

	Day 1	Day 2	Day 3
Patient 1	A	B	C
Patient 2	D	E	F

Wide

Long

	Day	Value
Patient 1	Day 1	A
Patient 1	Day 2	B
Patient 1	Day 3	C
Patient 2	Day 1	D
Patient 2	Day 2	E
Patient 2	Day 3	F

CC-BY jhudatascience.org

What is wide/long data?

Data is stored *differently* in the tibble.

Here's a small dataset looking at vaccination rates over three months in Alabama.

Wide: has many columns

```
# A tibble: 1 × 4
  State    June_vacc_rate May_vacc_rate April_vacc_rate
  <chr>        <dbl>       <dbl>          <dbl>
1 Alabama      0.516       0.514         0.511
```

Long: column names become data

```
# A tibble: 3 × 3
  State   name           value
  <chr>   <chr>         <dbl>
1 Alabama June_vacc_rate 0.516
2 Alabama May_vacc_rate  0.514
3 Alabama April_vacc_rate 0.511
```

What is wide/long data?

Wide: multiple columns per individual, values spread across multiple columns

```
# A tibble: 2 × 4
  State    June_vacc_rate May_vacc_rate April_vacc_rate
  <chr>        <dbl>       <dbl>          <dbl>
1 Alabama      0.516       0.514         0.511
2 Alaska       0.627       0.626         0.623
```

Long: multiple rows per observation, a single column contains the values

```
# A tibble: 6 × 3
  State   name       value
  <chr>  <chr>     <dbl>
1 Alabama June_vacc_rate 0.516
2 Alabama May_vacc_rate  0.514
3 Alabama April_vacc_rate 0.511
4 Alaska  June_vacc_rate 0.627
5 Alaska  May_vacc_rate  0.626
6 Alaska  April_vacc_rate 0.623
```

What is wide/long data?

<https://github.com/gadenbuie/tidyexplain/blob/main/images/tidyr-pivoting.gif>

wide

id	x	y	z
1	a	c	e
2	b	d	f

Why do we need to switch between wide/long data?

Wide: Easier for humans to read

```
# A tibble: 2 × 4
  State    June_vacc_rate May_vacc_rate April_vacc_rate
  <chr>        <dbl>       <dbl>          <dbl>
1 Alabama      0.516       0.514         0.511
2 Alaska       0.627       0.626         0.623
```

Long: Easier for R to make plots & do analysis

```
# A tibble: 6 × 3
  State   name        value
  <chr>   <chr>      <dbl>
1 Alabama June_vacc_rate 0.516
2 Alabama May_vacc_rate  0.514
3 Alabama April_vacc_rate 0.511
4 Alaska  June_vacc_rate 0.627
5 Alaska  May_vacc_rate  0.626
6 Alaska  April_vacc_rate 0.623
```

Pivoting using the **tidyr** package (part of **tidyverse**)

We will be talking about:

- `pivot_longer` - make multiple columns into variables, (wide to long)
- `pivot_wider` - make a variable into multiple columns, (long to wide)

The `reshape` command exists. Its arguments are considered more confusing, so we don't recommend it.

You might see old functions `gather` and `spread` when googling. These are older iterations of `pivot_longer` and `pivot_wider`, respectively.

pivot_longer...

Reshaping data from wide to long

`pivot_longer()` - puts column data into rows (`tidyverse` package)

- First describe which columns we want to “pivot_longer”

```
{long_data} <- {wide_data} |> pivot_longer(cols = {columns to pivot})
```

Reshaping data from wide to long

ex_wide

```
# A tibble: 2 × 4
  State    June_vacc_rate May_vacc_rate April_vacc_rate
  <chr>        <dbl>       <dbl>          <dbl>
1 Alabama      0.516       0.514         0.511
2 Alaska       0.627       0.626         0.623
```

```
ex_long <- ex_wide |> pivot_longer(cols = ends_with("rate"))
ex_long
```

```
# A tibble: 6 × 3
  State   name           value
  <chr>   <chr>         <dbl>
1 Alabama June_vacc_rate 0.516
2 Alabama May_vacc_rate  0.514
3 Alabama April_vacc_rate 0.511
4 Alaska  June_vacc_rate 0.627
5 Alaska  May_vacc_rate  0.626
6 Alaska  April_vacc_rate 0.623
```

GUT CHECK!

What does `pivot_longer()` do?

- A. Summarize data
- B. Import data
- C. Reshape data

Reshaping wide to long: Better column names

`pivot_longer()` - puts column data into rows (`tidyverse` package)

- First describe which columns we want to “pivot_longer”
- `names_to` = new name for old columns
- `values_to` = new name for old cell values

```
{long_data} <- {wide_data} |> pivot_longer(cols = {columns to pivot},  
                                         names_to = {name for old columns},  
                                         values_to = {name for cell values})
```

Reshaping wide to long: Better column names

Newly created column names (“Month” and “Rate”) are enclosed in quotation marks. It helps us be more specific than “name” and “value”.

```
ex_long <- ex_wide |> pivot_longer(cols = ends_with("rate"),
                                         names_to = "Month",
                                         values_to = "Rate")
```

```
ex_long
```

```
# A tibble: 6 × 3
  State    Month        Rate
  <chr>   <chr>       <dbl>
1 Alabama June_vacc_rate 0.516
2 Alabama May_vacc_rate  0.514
3 Alabama April_vacc_rate 0.511
4 Alaska   June_vacc_rate 0.627
5 Alaska   May_vacc_rate  0.626
6 Alaska   April_vacc_rate 0.623
```

Data used: Nitrate exposure

Let's look at some data on levels of nitrate in water from Washington. This dataset reports the amount of people in Washington exposed to excess levels of nitrate in their water between 1999 and 2020.

```
wide_nitrate <-  
  read_csv(file = "https://daseh.org/data/Nitrate_Exposure_for_WA_Public_Water_Systems_byquarter_data.csv")  
head(wide_nitrate)  
  
# A tibble: 6 × 11  
#>   year quarter pop_on_sampled_PWS `pop_0-3ug/L` `pop_>3-5ug/L` `pop_>5-10ug/L`  
#>   <dbl> <chr>       <dbl>        <dbl>        <dbl>        <dbl>  
#> 1 1999 Q1         106720      67775        0        32  
#> 2 1999 Q2          85541      55476        0       212  
#> 3 1999 Q3         559137     319252     231186       212  
#> 4 1999 Q4          26995      25969       420        0  
#> 5 2000 Q1          34793      5904        0        92  
#> 6 2000 Q2         184521     157396        0        32  
#> # 5 more variables: `pop_>10-20ug/L` <dbl>, `pop_>20ug/L` <dbl>,  
#> # `pop_on_PWS_with_non-detect` <dbl>, pop_exposed_to_exceedances <dbl>,  
#> # perc_pop_exposed_to_exceedances <dbl>
```

Mission: Average population exposed by concentration

Let's imagine we want to see what proportion of the population was exposed to different nitrate concentrations. Results should look something like:

```
# A tibble: 3 × 2
  concentration_cat avg_prop_exposedpop
  <chr>                <dbl>
1 0-3ug/L              0.593
2 10-20ug/L            0.000678
3 more_than_20ug/L     0.000129
```

Remove some columns we don't need

```
wide_nitrate <- wide_nitrate |>
  select(!ends_with("exceedances"))
wide_nitrate

# A tibble: 88 × 9
  year quarter pop_on_sampled_PWS `pop_0-3ug/L` `pop_>3-5ug/L` `pop_>5-10ug/L` 
  <dbl> <chr>           <dbl>          <dbl>          <dbl>          <dbl> 
1 1999 Q1            106720        67775          0         32
2 1999 Q2            85541         55476          0        212
3 1999 Q3            559137        319252        231186        212
4 1999 Q4            26995         25969         420          0
5 2000 Q1            34793         5904          0         92
6 2000 Q2            184521        157396          0         32
7 2000 Q3            42081         20407         345          0
8 2000 Q4            407219        358828        995        412
9 2001 Q1            90054         49552         150          0
10 2001 Q2            83521         43633        2536        90
# ... 78 more rows
# ... 3 more variables: `pop_>10-20ug/L` <dbl>, `pop_>20ug/L` <dbl>,
#   `pop_on_PWS_with_non-detect` <dbl>
```

Reshaping data from wide to long

```
long_nitrate <- wide_nitrate |>  
  pivot_longer(!c(year, quarter, pop_on_sampled_PWS))  
long_nitrate
```

A tibble: 528 × 5

	year	quarter	pop_on_sampled_PWS	name	value
	<dbl>	<chr>		<dbl>	<dbl>
1	1999	Q1	106720	pop_0-3ug/L	67775
2	1999	Q1	106720	pop_>3-5ug/L	0
3	1999	Q1	106720	pop_>5-10ug/L	32
4	1999	Q1	106720	pop_>10-20ug/L	0
5	1999	Q1	106720	pop_>20ug/L	0
6	1999	Q1	106720	pop_on_PWS_with_non-detect	38913
7	1999	Q2	85541	pop_0-3ug/L	55476
8	1999	Q2	85541	pop_>3-5ug/L	0
9	1999	Q2	85541	pop_>5-10ug/L	212
10	1999	Q2	85541	pop_>10-20ug/L	60
# 518 more rows					

Reshaping data from wide to long

Un-pivoted columns (year, quarter, pop_on_sampled_PWS) are still columns.

long_nitrate

```
# A tibble: 528 × 5
  year quarter pop_on_sampled_PWS name          value
  <dbl> <chr>   <dbl> <chr>        <dbl>
1 1999  Q1     106720  pop_0-3ug/L    67775
2 1999  Q1     106720  pop_>3-5ug/L    0
3 1999  Q1     106720  pop_>5-10ug/L   32
4 1999  Q1     106720  pop_>10-20ug/L   0
5 1999  Q1     106720  pop_>20ug/L    0
6 1999  Q1     106720  pop_on_PWS_with_non-detect 38913
7 1999  Q2     85541   pop_0-3ug/L    55476
8 1999  Q2     85541   pop_>3-5ug/L    0
9 1999  Q2     85541   pop_>5-10ug/L   212
10 1999  Q2    85541   pop_>10-20ug/L   60
# ... 518 more rows
```

Cleaning up long data

Let's make the conc_count into a proportion.

```
long_nitrate <- long_nitrate |>
  mutate(conc_prop = value / pop_on_sampled_PWS)
long_nitrate

# A tibble: 528 × 6
  year quarter pop_on_sampled_PWS name          value conc_prop
  <dbl> <chr>      <dbl> <chr>        <dbl>     <dbl>
1 1999 Q1       106720 pop_0-3ug/L 67775   0.635
2 1999 Q1       106720 pop_>3-5ug/L 0       0
3 1999 Q1       106720 pop_>5-10ug/L 32      0.000300
4 1999 Q1       106720 pop_>10-20ug/L 0       0
5 1999 Q1       106720 pop_>20ug/L 0       0
6 1999 Q1       106720 pop_on_PWS_with_non-detect 38913   0.365
7 1999 Q2       85541  pop_0-3ug/L 55476   0.649
8 1999 Q2       85541  pop_>3-5ug/L 0       0
9 1999 Q2       85541  pop_>5-10ug/L 212    0.00248
10 1999 Q2      85541  pop_>10-20ug/L 60     0.000701
# ... 518 more rows
```

Mission: Average population exposed by concentration

Now our data is more tidy, and we can take the averages easily!

```
long_nitrate |>  
  group_by(name) |>  
  summarize("avg_prop_exposedpop" = mean(conc_prop))  
  
# A tibble: 6 × 2  
  name          avg_prop_exposedpop  
  <chr>           <dbl>  
1 pop_0-3ug/L      0.593  
2 pop_>10-20ug/L    0.000678  
3 pop_>20ug/L       0.000129  
4 pop_>3-5ug/L        0.182  
5 pop_>5-10ug/L       0.0189  
6 pop_on_PWS_with_non-detect 0.206
```

Reshaping data from wide to long

There are many ways to **select** the columns we want. Check out https://dplyr.tidyverse.org/reference/dplyr_tidy_select.html to look at more column selection options.

pivot_wider...

Reshaping data from long to wide

`pivot_wider()` - spreads row data into columns (`tidyverse` package)

- `names_from` = the old column whose contents will be spread into multiple new column names.
- `values_from` = the old column whose contents will fill in the values of those new columns.

```
{wide_data} <- {long_data} |>  
  pivot_wider(names_from = {Old column name: contains new column names},  
             values_from = {Old column name: contains new cell values})
```

Reshaping data from long to wide

We can use `pivot_wider` to convert long data to wide format. Let's try it with the vaccine data from earlier.

ex_long

```
# A tibble: 6 × 3
  State    Month        Rate
  <chr>   <chr>      <dbl>
1 Alabama June_vacc_rate 0.516
2 Alabama May_vacc_rate  0.514
3 Alabama April_vacc_rate 0.511
4 Alaska  June_vacc_rate 0.627
5 Alaska  May_vacc_rate  0.626
6 Alaska  April_vacc_rate 0.623
```

Reshaping data from long to wide

We can use `pivot_wider` to convert long data to wide format. Let's try it with the vaccine data from earlier.

```
ex_wide2 <- ex_long |> pivot_wider(names_from = "Month",
                                         values_from = "Rate")  
ex_wide2
```

```
# A tibble: 2 × 4  
  State    June_vacc_rate May_vacc_rate April_vacc_rate  
  <chr>        <dbl>       <dbl>          <dbl>  
1 Alabama      0.516       0.514          0.511  
2 Alaska       0.627       0.626          0.623
```

Reshaping nitrate exposure data

Let's go back to the nitrate exposure dataset. What if we wanted to make a wide version of the data that displayed the proportion of people at each level of nitrate exposure, with each quarter as a column?

```
long_nitrate
```

```
# A tibble: 528 × 6
```

	year	quarter	pop_on_sampled_PWS	name	value	conc_prop
	<dbl>	<chr>	<dbl>	<chr>	<dbl>	<dbl>
1	1999	Q1	106720	pop_0-3ug/L	67775	0.635
2	1999	Q1	106720	pop_>3-5ug/L	0	0
3	1999	Q1	106720	pop_>5-10ug/L	32	0.000300
4	1999	Q1	106720	pop_>10-20ug/L	0	0
5	1999	Q1	106720	pop_>20ug/L	0	0
6	1999	Q1	106720	pop_on_PWS_with_non-detect	38913	0.365
7	1999	Q2	85541	pop_0-3ug/L	55476	0.649
8	1999	Q2	85541	pop_>3-5ug/L	0	0
9	1999	Q2	85541	pop_>5-10ug/L	212	0.00248
10	1999	Q2	85541	pop_>10-20ug/L	60	0.000701
# 518 more rows						

Reshaping nitrate exposure data

Drop some columns we don't need.

```
long_nitrate <- long_nitrate |>
  select(!c(pop_on_sampled_PWS, value))
long_nitrate

# A tibble: 528 × 4
  year quarter name          conc_prop
  <dbl> <chr>   <chr>        <dbl>
1 1999  Q1     pop_0-3ug/L  0.635
2 1999  Q1     pop_>3-5ug/L 0
3 1999  Q1     pop_>5-10ug/L 0.000300
4 1999  Q1     pop_>10-20ug/L 0
5 1999  Q1     pop_>20ug/L  0
6 1999  Q1     pop_on_PWS_with_non-detect 0.365
7 1999  Q2     pop_0-3ug/L  0.649
8 1999  Q2     pop_>3-5ug/L 0
9 1999  Q2     pop_>5-10ug/L 0.00248
10 1999  Q2    pop_>10-20ug/L 0.000701
# ... 518 more rows
```

Reshaping nitrate exposure data

Pivot the data!

```
wide_nitrate <- long_nitrate |>
  pivot_wider(names_from = "quarter", values_from = "conc_prop")
wide_nitrate

# A tibble: 132 × 6
  year    name      Q1     Q2     Q3     Q4
  <dbl>   <chr>   <dbl>   <dbl>   <dbl>   <dbl>
1 1999 pop_0-3ug/L 0.635   0.649   0.571   0.962
2 1999 pop_>3-5ug/L 0        0        0.413   0.0156
3 1999 pop_>5-10ug/L 0.000300 0.00248  0.000379 0
4 1999 pop_>10-20ug/L 0       0.000701 0       0
5 1999 pop_>20ug/L 0        0        0       0
6 1999 pop_on_PWS_with_non-detect 0.365   0.348   0.0152  0.0224
7 2000 pop_0-3ug/L 0.170   0.853   0.485   0.881
8 2000 pop_>3-5ug/L 0        0        0.00820 0.00244
9 2000 pop_>5-10ug/L 0.00264  0.000173 0       0.00101
10 2000 pop_>10-20ug/L 0       0        0       0
# ... 122 more rows
```

Summary

- `tidyverse` package helps us convert between wide and long data
- `pivot_longer()` goes from wide -> long
 - Specify columns you want to pivot
 - Specify `names_to =` and `values_to =` for custom naming
- `pivot_wider()` goes from long -> wide
 - Specify `names_from =` and `values_from =`

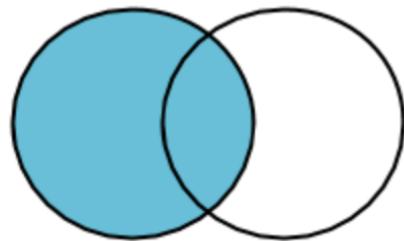
Lab Part 1

□ [Class Website](#)

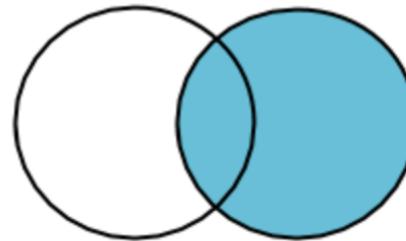
□ [Lab](#)

Joining

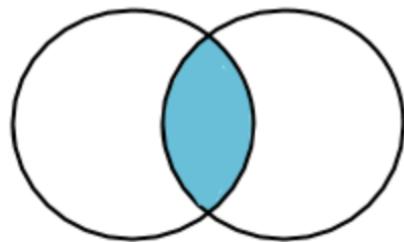
“Combining datasets”



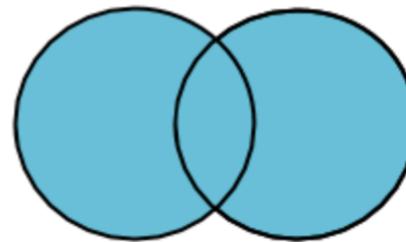
Left Join



Right Join



Inner Join



**Full Outer
Join**

Joining in **dplyr**

- Merging/joining data sets together - usually on key variables, usually "id"
- `?join` - see different types of joining for **dplyr**
- `inner_join(x, y)` - only rows that match for x and y are kept
- `full_join(x, y)` - all rows of x and y are kept
- `left_join(x, y)` - all rows of x are kept even if not merged with y
- `right_join(x, y)` - all rows of y are kept even if not merged with x
- `anti_join(x, y)` - all rows from x not in y keeping just columns from x.

Merging: Simple Data

```
data_As <- read_csv(  
  file = "https://daseh.org/data/data_As_1.csv")  
data_cold <- read_csv(  
  file = "https://daseh.org/data/data_cold_1.csv")
```

data_As

```
# A tibble: 2 × 3  
  State    June_vacc_rate May_vacc_rate  
  <chr>        <dbl>          <dbl>  
1 Alabama      0.516         0.514  
2 Alaska       0.627         0.626
```

data_cold

```
# A tibble: 2 × 2  
  State    April_vacc_rate  
  <chr>        <dbl>  
1 Maine      0.795  
2 Alaska     0.623
```

Inner Join

<https://github.com/gadenbuie/tidyexplain/blob/main/images/inner-join.gif>

inner_join(x, y)	
1	x1
2	x2
3	x3
1	y1
2	y2
4	y4

Inner Join

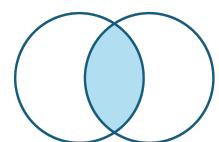
```
ij <- inner_join(data_As, data_cold)
```

```
Joining with `by = join_by(State)`
```

```
ij
```

```
# A tibble: 1 × 4
```

```
  State June_vacc_rate May_vacc_rate April_vacc_rate
  <chr>      <dbl>        <dbl>          <dbl>
1 Alaska     0.627       0.626        0.623
```



Left Join

<https://raw.githubusercontent.com/gadenbuie/tidyexplain/main/images/left-join.gif>

`left_join(x, y)`

1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

Left Join

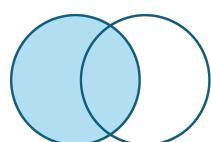
“Everything to the left of the comma”

```
lj <- left_join(data_As, data_cold)
```

Joining with `by = join_by(State)`

```
lj
```

```
# A tibble: 2 × 4
  State    June_vacc_rate May_vacc_rate April_vacc_rate
  <chr>        <dbl>       <dbl>          <dbl>
1 Alabama     0.516       0.514          NA
2 Alaska      0.627       0.626         0.623
```

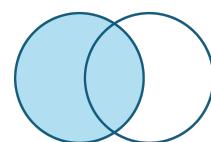


Install **tidylog** package to log outputs

```
# install.packages("tidylog")
library(tidylog)
left_join(data_As, data_cold)

Joining with `by = join_by(State)`
left_join: added one column (April_vacc_rate)
> rows only in data_As 1
> rows only in data_cold (1)
> matched rows 1
> ===
> rows total 2

# A tibble: 2 × 4
  State    June_vacc_rate  May_vacc_rate April_vacc_rate
  <chr>        <dbl>        <dbl>            <dbl>
1 Alabama      0.516       0.514            NA
2 Alaska       0.627       0.626            0.623
```



Right Join

<https://raw.githubusercontent.com/gadenbuie/tidyexplain/main/images/right-join.gif>

right_join(x, y)			
1	x1	1	y1
2	x2	2	y2
3	x3	4	y4

Right Join

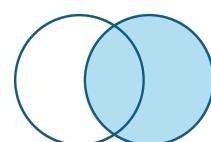
“Everything to the right of the comma”

```
rj <- right_join(data_As, data_cold)
```

```
Joining with `by = join_by(State)`  
right_join: added one column (April_vacc_rate)  
> rows only in data_As (1)  
> rows only in data_cold 1  
> matched rows 1  
> ===  
> rows total 2
```

```
rj
```

```
# A tibble: 2 × 4  
  State    June_vacc_rate  May_vacc_rate April_vacc_rate  
  <chr>        <dbl>        <dbl>        <dbl>  
1 Alaska      0.627       0.626       0.623  
2 Maine       NA          NA          0.795
```



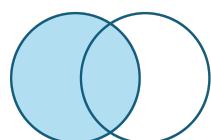
Left Join: Switching arguments

```
lj2 <- left_join(data_cold, data_As)
```

```
Joining with `by = join_by(State)`
left_join: added 2 columns (June_vacc_rate, May_vacc_rate)
> rows only in data_cold 1
> rows only in data_As (1)
> matched rows 1
> ===
> rows total 2
```

```
lj2
```

```
# A tibble: 2 × 4
  State April_vacc_rate June_vacc_rate May_vacc_rate
  <chr>      <dbl>          <dbl>          <dbl>
1 Maine       0.795         NA             NA
2 Alaska      0.623         0.627         0.626
```



Full Join

<https://raw.githubusercontent.com/gadenbuie/tidyexplain/main/images/full-join.gif>

`full_join(x, y)`

1	x1
2	x2
3	x3

1	y1
2	y2
4	y4

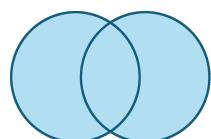
Full Join

```
fj <- full_join(data_As, data_cold)
```

```
Joining with `by = join_by(State)`
full_join: added one column (April_vacc_rate)
> rows only in data_As 1
> rows only in data_cold 1
> matched rows 1
> ===
> rows total 3
```

```
fj
```

```
# A tibble: 3 × 4
  State    June_vacc_rate May_vacc_rate April_vacc_rate
  <chr>        <dbl>       <dbl>          <dbl>
1 Alabama      0.516       0.514          NA
2 Alaska       0.627       0.626         0.623
3 Maine        NA          NA            0.795
```



"includes duplicates"

```
data_As <- read_csv(  
  file = "https://daseh.org/data/data_As_2.csv")  
data_cold <- read_csv(  
  file = "https://daseh.org/data/data_cold_2.csv")
```

data_As

```
# A tibble: 2 × 2  
  State    state_bird  
  <chr>    <chr>  
1 Alabama  wild turkey  
2 Alaska   willow ptarmigan
```

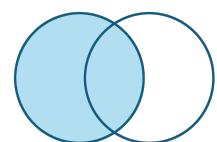
data_cold

```
# A tibble: 3 × 3  
  State    vacc_rate month  
  <chr>    <dbl>  <chr>  
1 Maine     0.795 April  
2 Alaska    0.623 April  
3 Alaska    0.626 May
```

“includes duplicates”

```
lj <- left_join(data_As, data_cold)

Joining with `by = join_by(State)`
left_join: added 2 columns (vacc_rate, month)
> rows only in data_As 1
> rows only in data_cold (1)
> matched rows 2 (includes duplicates)
> ===
> rows total 3
```



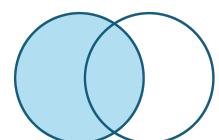
“includes duplicates”

Data including the joining column (“State”) has been duplicated.

1j

```
# A tibble: 3 × 4
  State    state_bird      vacc_rate month
  <chr>   <chr>           <dbl>     <chr>
1 Alabama wild turkey      NA        <NA>
2 Alaska  willow ptarmigan 0.623     April
3 Alaska  willow ptarmigan 0.626     May
```

Note that “Alaska willow ptarmigan” appears twice.



“includes duplicates”

<https://github.com/gadenbuie/tidyexplain/blob/main/images/left-join-extra.gif>

`left_join(x, y)`

1	x1
2	x2
3	x3

1	y1
2	y2
4	y4
2	y5

Stop **tidylog**

`unloadNamespace()` does the opposite of `library()`.

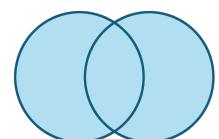
```
unloadNamespace("tidylog")
```

Using the `by` argument

By default joins use the intersection of column names. If `by` is specified, it uses that.

```
full_join(data_As, data_cold, by = "State")
```

```
# A tibble: 4 × 4
  State    state_bird      vacc_rate month
  <chr>   <chr>           <dbl>     <chr>
1 Alabama wild turkey     NA        <NA>
2 Alaska  willow ptarmigan 0.623    April
3 Alaska  willow ptarmigan 0.626    May
4 Maine   <NA>            0.795    April
```

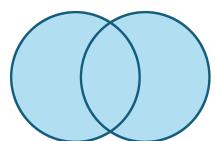


Using the **by** argument

You can join based on multiple columns by using something like `by = c(col1, col2)`.

If the datasets have two different names for the same data, use:

```
full_join(x, y, by = c("a" = "b"))
```



anti_join: what's missing

Entries in data_As but not in data_cold

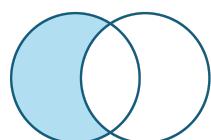
```
anti_join(data_As, data_cold, by = "State")
```

```
# A tibble: 1 × 2
  State    state_bird
  <chr>    <chr>
1 Alabama wild turkey
```

Entries in data_cold but not in data_As

```
anti_join(data_cold, data_As, by = "State") # order switched
```

```
# A tibble: 1 × 3
  State    vacc_rate month
  <chr>      <dbl>   <chr>
1 Maine      0.795 April
```



GUT CHECK!

Why use `join` functions?

- A. Combine different data sources
- B. Connect Rmd to other files
- C. Using one data source is too easy and we want our analysis ~ fancy ~

Summary

- Merging/joining data sets together - assumes all column names that overlap
 - use the `by = c("a" = "b")` if they differ
- `inner_join(x, y)` - only rows that match for x and y are kept
- `full_join(x, y)` - all rows of x and y are kept
- `left_join(x, y)` - all rows of x are kept even if not merged with y
- `right_join(x, y)` - all rows of y are kept even if not merged with x
- Use the `tidylog` package for a detailed summary
- `anti_join(x, y)` shows what is only in x (missing from y)

Lab Part 2

- [Class Website](#)
- [Lab](#)
- [Day 6 Cheatsheet](#)
- [Posit's tidyverse Cheatsheet](#)
- [Posit's dplyr Cheatsheet](#)



Image by [Gerd Altmann from Pixabay](#)

Additional Slides

Getting the set difference with **setdiff**

We might want to determine what indexes ARE in the first dataset that AREN'T in the second.

For this to work, the datasets need the same columns.

We'll just select the index using `select()`.

```
A_states <- data_A %> select(State)  
cold_states <- data_cold %> select(State)
```

Getting the set difference with **setdiff**

States in A_states but not in cold_states

```
dplyr::setdiff(A_states, cold_states)
```

```
# A tibble: 1 × 1
  State
  <chr>
1 Alabama
```

States in cold_states but not in A_states

```
dplyr::setdiff(cold_states, A_states)
```

```
# A tibble: 1 × 1
  State
  <chr>
1 Maine
```

Getting the set difference with **setdiff**

Why did we use `dplyr::setdiff`?

There is a base R function, also called `setdiff` that requires vectors.

In other words, we use `dplyr::` to be specific about the package we want to use.

More set operations can be found here:

<https://dplyr.tidyverse.org/reference/setops.html>

Fast manipulation using **collapse** package

<https://sebkrantz.github.io/collapse/>

Might be helpful if your data is very large. However, **dplyr** and **tidyverse** functions are great for most applications.